

Krishna vamsi Dhulipalla

(+1) 540-558-3528 | dhulipallakrishnavamsi@gmail.com | [GitHub](#) | [LinkedIn](#) | [Portfolio](#) | [Personal Chatbot](#)

Experience

Software Engineer - AI Platform Cloud Systems LLC US Remote	Jul 2025 - Present
<ul style="list-style-type: none">Architected an agentic workflow using LangGraph and ReAct prompting to automate SQL generation; automated ~65% of ad-hoc data-auditing requests from internal stakeholders, reducing the average response time from 4 hrs to <2 minOptimized data ingestion performance by rebuilding ETL pipelines with batched I/O, incremental refresh logic, and dependency pruning, cutting daily execution runtime by 25%Improved production reliability by shipping the agent service on Kubernetes with autoscaling and rolling deploys, then adding alerts and rollback steps for failed releasesImproved cross-service reliability by implementing Pydantic schema validation and contract tests, catching 100% of integration bugs during CI before they reached production	
Machine Learning Engineer Virginia Tech, Dept. of Plant Sciences Blacksburg, VA	Aug 2024 - Jul 2025
<ul style="list-style-type: none">Increased genomics sequence classification throughput by 32% by applying LoRA and soft prompting; packaged repeatable PyTorch pipelines that cut per-experiment training time by 4.5 hoursEngineered a distributed pipeline for 100+ multi-GPU experiments on HPC clusters; implemented custom checkpoint-resume logic that prevented data loss during preemptive node shutdowns, reducing total compute waste by ~15%Reduced research environment setup time from hours to minutes by containerizing fine-tuned models with Docker and managing the experimental lifecycle (versions, hyperparameters, and weights) via MLflow	
Software Engineer UJR Technologies Pvt Ltd Hyderabad, India	Jul 2021 - Dec 2022
<ul style="list-style-type: none">Developed standardized REST APIs and UI workflows using FastAPI and JSON Schema, reducing cross-service integration defects by 40% across 3 core product modulesEngineered model-serving endpoints with automated input validation and deployment health checks, lowering prediction-related failures by 30% for ML-driven featuresAutomated CI/CD pipelines via GitHub Actions with comprehensive test coverage and scripted rollback procedures, decreasing release failures by 20% across production environments	

Projects

Autonomous Multi-Agent Web UI Automation System

- Developed a multi-agent system using LangGraph and Playwright to navigate non-deterministic UI changes across 5 high-complexity SaaS platforms
- Increased task completion success rate from 68% to 94% by implementing a two-stage verification loop with step-level assertions and exponential backoff for dynamic DOM states
- Integrated LangSmith traces for observability, reducing the mean-time-to-debug for broken selectors by 14 min per incident

Proxy TuNER: Advancing Cross-Domain Named Entity Recognition through Proxy Tuning

- Improved cross-domain NER F1-score by 8% by implementing a proxy-tuning approach for LLaMA 2 models (7B, 7B-Chat, 13B) using logit ensembling and gradient reversal
- Optimized inference performance by 30% and reduced training costs by 70% through distributed execution and model-path optimizations in PyTorch

IntelliMeet: AI-Enabled Decentralized Video Conferencing App

- Architected a secure, decentralized video platform using WebRTC and federated learning to maintain data privacy while sharing only aggregated model updates
- Reduced call dropouts by 25% by engineering network recovery logic and on-device RetinaFace attention detection for client-side quality adaptation

SKILLS

Languages: Python, SQL, TypeScript, JavaScript, MongoDB

ML & AI Frameworks: PyTorch, Transformers, LangChain, LangGraph, LoRA, RAG, NLP

Data & Infrastructure: Docker, Kubernetes, Apache Airflow, MLflow, Redis, FAISS, AWS, GCP, Git

Tools & Observability: LangSmith, Grafana, CI/CD (GitHub Actions/Jenkins), Weights & Biases, Linux

Publications

- Predicting Circadian Transcription in mRNAs and lncRNAs, IEEE BIBM 2024
- DNA Foundation Models for Cross-Species TF Binding Prediction, NeurIPS ML in CompBio 2025
- Multi-omics atlas of the plant nuclear envelope, Science Advances (under review) 2025, University of California, Berkeley

Education

Virginia Tech | Master of Science, Computer Science | Dec 2024

GPA 3.9/4.0

Vel Tech University | Bachelor of Technology, Computer Science and Engineering

GPA 8.24/10