

Krishna vamsi Dhulipalla

(+1) 540-558-3528 | dhulipallakrishnavamsi@gmail.com | GitHub: [Krishna-dhulipalla](#) | LinkedIn: [krishnavamsidhulipalla](#)
[Portfolio](#) | [Personal Chatbot](#)

Summary

Software and ML Engineer with 3+ years of experience building agentic LLM systems, multi-agent workflows, and AI-driven internal tools. Skilled in Python, distributed training, and cloud-native development, model evaluation, and scalable deployment. Experienced in designing RAG pipelines, automating data workflows, and delivering reliable AI services used by cross-functional teams.

Experience

Software Engineer – AI/ML Infrastructure

Cloud Systems LLC

Jul. 2025 - Present

Remote

- Designed multi-agent workflows with planning, tool-use, and schema-verified retrieval, enabling reliable automation of complex SQL + API tasks and reducing manual review cycles by 40%
- Reduced daily ETL runtime by 25% by rebuilding SQL and model-evaluation workflows with batched I/O, pruned steps, and incremental loads, keeping agent data up to date
- Ensured service reliability and uptime (>99.9%) by containerizing agent services on Kubernetes with rolling updates and horizontal pod autoscaling
- Shortened model integration delays by 20% by defining API contracts early and adding CI schema validation to speed up feature delivery

Machine Learning Engineer

Virginia Tech

Aug. 2024 - Jul. 2025

Blacksburg, VA

- Improved genomic sequence classification throughput by 30% by fine-tuning LLMs with LoRA and soft prompting and packaging training/eval as repeatable runs
- Shortened multi-GPU training cycles from days to hours by orchestrating 100+ biosequence experiments with scheduling and checkpoint reuse on HPC clusters
- Led development of an AI search interface with LangChain and semantic retrieval, replacing manual literature review for research teams
- Reduced setup time for researchers by 25% by containerizing fine-tuned models with MLflow for reproducible deployment

Data Engineer

UJR Technologies Pvt Ltd

Jul. 2021 - Dec. 2022

Hyderabad, India

- Lowered analytics latency by 30% by moving ETL to a Kafka+Spark streaming pipeline for real-time financial reporting
- Reduced release rollbacks by 40% by shipping Dockerized microservices to AWS EKS with health checks and staged rollouts
- Boosted dashboard query speed by 40% by redesigning Snowflake schemas and materialized views for common reporting paths
- Maintained 99.9% uptime by adding CloudWatch metrics and alerts for ingestion lag, error rates, and throughput
- Cut client report generation time by 35% by integrating Power BI with Snowflake models co-designed with the BI and frontend teams

Skills

Programming:

Python, R, SQL, JavaScript, TypeScript, Node.js, MongoDB, FastAPI, Go

ML & AI Frameworks:

PyTorch, TensorFlow, Scikit-learn, JAX, Hugging Face Transformers, OpenCV

Multi-Agent & LLM:

LangChain, LangGraph, AutoGen, MCP, LLM Fine-Tuning, RAG, OpenAI SDK, Prompt Engineering, evaluation harnesses, guardrails, Llmaindex

Data Engineering & Cloud:

Apache Spark, Kafka, Airflow, AWS (S3, Redshift, ECS, SageMaker), GCP (GCS, BigQuery, Dataflow), Snowflake, Vector Databases (FAISS, Pinecone), Redis

MLOps & Infrastructure:

Docker, Kubernetes, Jenkins, MLflow, CI/CD, Weights & Biases, Git, Linux, GitHub Actions, CloudWatch, Grafana

Other:

LangSmith, Evals, Experiment Tracking, A/B Testing, Model Evaluation, Data Visualization, Streamlit, Semantic Search, PowerBI

Education

Virginia Tech

M.S. in Computer Science

Blacksburg, VA

CGPA - 3.9/4

Vel Tech University

Bachelor's in computer science

Chennai, India

CGPA - 8.24/10

Projects

Multi-Agent AI System: Community & Hazard Intelligence Map

- Unified official feeds and user reports on a live map powered by LangGraph agent, lifting engagement 35%
- Built reporting, classification, and verification workflows to surface high-confidence alerts faster
- Fixed concurrency and map-overlay bugs with request guards and optimistic updates, improving update consistency
- Integrated MCP protocol for agent orchestration, aligning with emerging interoperability standards

Proxy TuNER: Advancing Cross-Domain Named Entity Recognition through Proxy Tuning

- Implemented a proxy-tuning approach for BERT using logit ensembling with domain-specific expert models, improving F1-score by 8% across diverse datasets
- Cut training cost 70% and sped up inference 30% through distributed runs and model-path optimizations
- Applied gradient reversal for domain-invariant feature learning, boosting cross-domain accuracy by up to 15%

IntelliMeet: AI-Enabled Decentralized Video Conferencing App

- Built a secure video conferencing platform with federated learning and encryption to protect user data
- Reduced meeting dropouts 25% by running on-device attention detection and handling network hiccups gracefully
- Delivered automatic notes with speech-to-text and summarization, removing manual minutes for every call
- Kept user data on device and encrypted traffic end-to-end; maintained 99.9% uptime with staged releases

Publications

Predicting Circadian Transcription in mRNAs and lncRNAs — IEEE BIBM 2024

Applied ML models to genomic data; improved transcription prediction accuracy, enabling deeper insights into biological rhythms. [DOI: [10.1109/BIBM62325.2024.10822684](https://doi.org/10.1109/BIBM62325.2024.10822684)]

DNA Foundation Models for Cross-Species TF Binding Prediction — NeurIPS ML in CompBio 2025

Leveraged DNABERT-style architectures for plant genomics; advanced cross-species binding prediction with improved generalization. [bioRxiv: [10.1101/2025.07.14.664780v1](https://doi.org/10.1101/2025.07.14.664780v1)]

Certifications

- NVIDIA – Building RAG Agents with LLMs (2025)
- NVIDIA – Deploying RAG Pipelines for Production at Scale (2025)
- AWS – Delivering Data-Driven Decisions (2024)
- Snowflake – End-to-End Data Engineering (2024)
- Google Cloud – Data Engineering Foundations (2024)
- AWS – Cloud Foundations (2022)